

Problem Statement:

A real estate agent wants help to predict the house price for regions in USA. He gave you the dataset to work on and we decided to use Linear Regression model. Creating a model that will help to estimate of what the house would sell for

```
In [1]: #importing Libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn import preprocessing, svm  
from sklearn.model_selection import train_test_split
```

```
In [2]: from sklearn.linear_model import LinearRegression
```

Data collection:

```
In [3]: #Reading the data
df=pd.read_csv(r"C:/Users/my pc/downloads/USA_Housing.csv")
df
```

Out[3]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Ad
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Fer 674\nLaurabu 3
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Suite 079\ Kathleen,
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Eliz Stravenue\ Danie WI 06
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\ FF
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\ AE (
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\ AP 30153
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 925\ 8489\nAPO AA 4
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy G Suite 076\ Joshua VI
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\ FF
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George F Apt. 509\ East N

5000 rows × 7 columns



In [4]: `df.head()` *#displays the first 5 rows*

Out[4]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry A 674\nLaurabury, N 3701
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson View Suite 079\nLal Kathleen, CA
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabe Stravenue\nDanieltow WI 06482
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO / 448;
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFF AE 093;

In [5]: `df.tail()` *#displays the last 5 rows*

Out[5]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991-3352
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2...

In [6]: `df.describe()`

Out[6]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                     5000 non-null   float64
1   Avg. Area House Age                  5000 non-null   float64
2   Avg. Area Number of Rooms            5000 non-null   float64
3   Avg. Area Number of Bedrooms         5000 non-null   float64
4   Area Population                      5000 non-null   float64
5   Price                               5000 non-null   float64
6   Address                             5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

In [8]: `df.isna().any()` *#checking for null values*

```
Out[8]: Avg. Area Income      False
Avg. Area House Age      False
Avg. Area Number of Rooms  False
Avg. Area Number of Bedrooms False
Area Population          False
Price                   False
Address                 False
dtype: bool
```

In [9]: `df.shape`

Out[9]: (5000, 7)

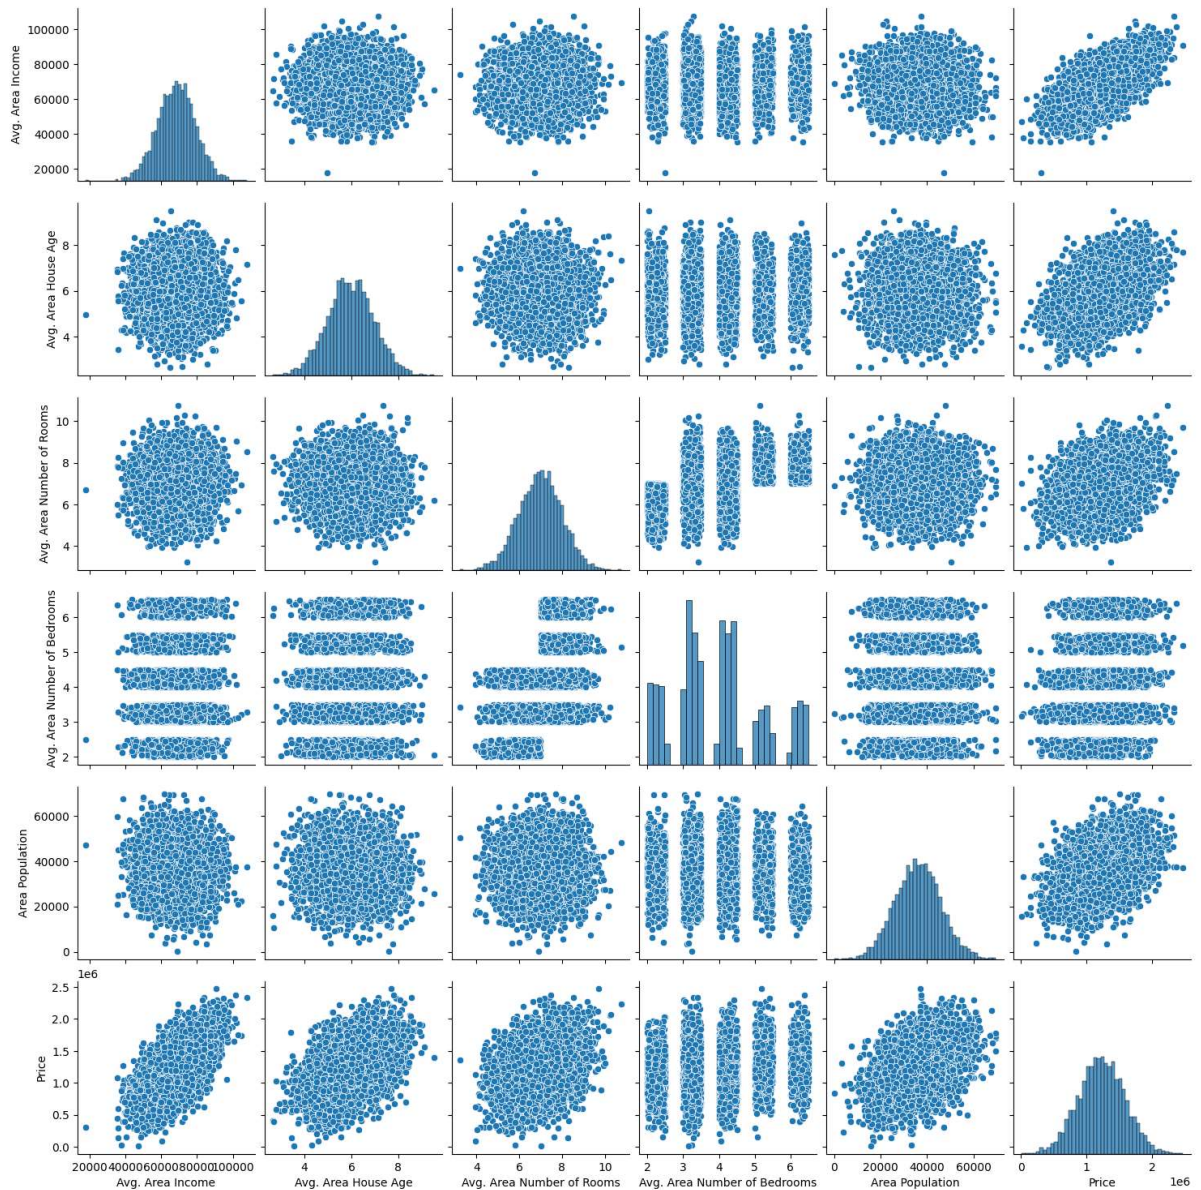
```
In [10]: df.columns
```

```
Out[10]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
              'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],  
              dtype='object')
```

Exploratory Data Analysis


```
In [11]: sns.pairplot(df)
```

```
Out[11]: <seaborn.axisgrid.PairGrid at 0x2767b386a70>
```




```
In [12]: sns.displot(df['Price'])
```

```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x2767aff6470>
```



```
In [13]: sns.displot(df['Area Population'])
```


```
Out[13]: <seaborn.axisgrid.FacetGrid at 0x27669beaaa0>
```



```
In [14]: Housedf=df[["Avg. Area Income","Avg. Area House Age","Avg. Area Number of Rooms",  
                    "Area Population","Price"]]
```

```
In [15]: sns.heatmap(Housedf.corr())
```

```
Out[15]: <Axes: >
```



To train the model

```
In [16]: x=Housedf[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
                    'Avg. Area Number of Bedrooms', 'Area Population']]
          y=df['Price']
```

```
In [17]: from sklearn.model_selection import train_test_split
          x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=42)
```

```
In [18]: from sklearn.linear_model import LinearRegression
```

```
In [19]: lm=LinearRegression()
          lm.fit(x_train,y_train)
```

```
Out[19]: LinearRegression
          LinearRegression()
```



```
In [20]: print(lm.intercept_)
```

```
-2644850.0694553573
```

```
In [21]: coeff_df=pd.DataFrame(lm.coef_,x.columns,columns=['coefficient'])  
coeff_df
```


```
Out[21]:
```

	coefficient
Avg. Area Income	21.664598
Avg. Area House Age	165789.776062
Avg. Area Number of Rooms	120587.850072
Avg. Area Number of Bedrooms	1431.988439
Area Population	15.248314

Evaluating model performance


```
In [23]: predictions=lm.predict(x_test)  
plt.scatter(y_test,predictions)
```

```
Out[23]: <matplotlib.collections.PathCollection at 0x2760120c3d0>
```



```
In [25]: sns.displot((y_test,predictions),bins=50)
```

```
Out[25]: <seaborn.axisgrid.FacetGrid at 0x27600bc3580>
```



```
In [40]: print("MAE:",metrics.mean_absolute_error(y_test,predictions))
print("MSE:",metrics.mean_squared_error(y_test,predictions))
print("RMSE:",np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

```
MAE: 81832.71748622792
MSE: 10401275405.858944
RMSE: 101986.64327184684
```

```
In [38]: from sklearn.metrics import r2_score
r2=r2_score(y_test,predictions)
print("R2 score:",r2)
```

```
R2 score: 0.918508746677805
```

conclusion:

It is best model ,because the r2 is high and the linear and normal distribution are best fitted for this dataset

